

# Customer Churn Analysis and Prediction for a Telecommunications Company



## Project Overview:

The project focuses on analyzing customer churn and building predictive models to understand the factors contributing to customer retention for a telecommunications company. Customer churn, which refers to customers discontinuing their services, is a crucial metric for telecom businesses, given that retaining existing customers is typically more cost-effective than acquiring new ones. To address this challenge, this project employs a comprehensive approach that includes feature engineering, encoding and transforming data, and the utilization of **Artificial Neural Networks (ANN)**.

The project leverages these techniques to analyze a dataset that provides insights into customer behavior and whether they have churned. By transforming and encoding the data, we make it suitable for ANN, a powerful machine learning tool for predictive modeling. ANN can learn complex patterns within the data and help identify factors contributing to customer churn.

## Outcome:

Based on the results, we can draw the following conclusions:

### Logistic Regression Model:

The logistic regression model achieved an accuracy of 80%, indicating that it correctly classified 80% of the samples. The model's performance is reasonably good, with higher precision and recall for Churn=No compared to Churn=Yes. The F1-score, which balances precision and recall, is also higher for Churn=No.

### OLS Regression Model:

The OLS regression model explains approximately 21.5% of the variance in the dependent variable (Churn). The model suggests that both tenure and monthly charges have statistically significant effects on churn. The coefficients provide insight into how these variables influence churn.

```
In [1]: import pandas as pd
import os
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
import statsmodels.api as sm
import statsmodels.api as smf
import statsmodels.stats.api as sms
```

```

from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import roc_curve, roc_auc_score, auc
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
from tensorflow import keras

warnings.filterwarnings("ignore")
%matplotlib inline

```

```

In [2]: directory = "F:\Jupyter\Telecom Churn Analysis"
filename = "data_telco_customer_churn.csv"

file_path = os.path.join(directory, filename)

if os.path.exists(file_path):
    data = pd.read_csv(file_path)
    print(data.head())

else:
    print("The file does not exists")

```

	Dependents	tenure	OnlineSecurity	OnlineBackup	\
0	Yes	9	No	No	
1	No	14	No	Yes	
2	No	64	Yes	No	
3	No	72	Yes	Yes	
4	No	3	No internet service	No internet service	

	InternetService	DeviceProtection	TechSupport	Contract	\
0	DSL	Yes	Yes	Month-to-month	
1	Fiber optic	Yes	No	Month-to-month	
2	DSL	Yes	Yes	Two year	
3	DSL	Yes	Yes	Two year	
4	No	No internet service	No internet service	Month-to-month	

	PaperlessBilling	MonthlyCharges	Churn
0	Yes	72.90	Yes
1	Yes	82.65	No
2	No	47.85	Yes
3	No	69.65	No
4	Yes	23.60	No

```

In [3]: data.sample(5)

```

```

Out[3]:

```

	Dependents	tenure	OnlineSecurity	OnlineBackup	InternetService	DeviceProtection	TechSupport	Contract	PaperlessBilling	MonthlyC
107	No	36	No	No	Fiber optic	No	No	Month-to-month	Yes	
4795	Yes	47	Yes	No	Fiber optic	Yes	No	Month-to-month	No	
335	No	65	No internet service	No internet service	No	No internet service	No internet service	Two year	Yes	
4843	No	46	Yes	No	DSL	Yes	Yes	Two year	No	
1324	No	4	Yes	No	DSL	Yes	Yes	Month-to-month	No	

```

In [4]: data.isnull().sum()

```

```

Out[4]:
Dependents      0
tenure          0
OnlineSecurity  0
OnlineBackup    0
InternetService 0
DeviceProtection 0
TechSupport     0
Contract        0
PaperlessBilling 0
MonthlyCharges  0
Churn           0
dtype: int64

```

```

In [5]: data.info

```

```
Out[5]: <bound method DataFrame.info of          Dependents  tenure      OnlineSecurity      OnlineBackup \
0          Yes          9          No          No
1          No         14          No          Yes
2          No         64          Yes          No
3          No         72          Yes          Yes
4          No          3  No internet service  No internet service
...          ...          ...          ...          ...
4925         No         15          No          No
4926        Yes         10          No          No
4927         No         58  No internet service  No internet service
4928         No          1          No          No
4929        Yes          4  No internet service  No internet service
```

```
          InternetService  DeviceProtection      TechSupport \
0          DSL          Yes          Yes
1      Fiber optic          Yes          No
2          DSL          Yes          Yes
3          DSL          Yes          Yes
4          No  No internet service  No internet service
...          ...          ...          ...
4925      Fiber optic          Yes          Yes
4926      Fiber optic          Yes          Yes
4927          No  No internet service  No internet service
4928      Fiber optic          No          No
4929          No  No internet service  No internet service
```

```
          Contract  PaperlessBilling  MonthlyCharges  Churn
0      Month-to-month          Yes          72.90      Yes
1      Month-to-month          Yes          82.65      No
2          Two year          No          47.85      Yes
3          Two year          No          69.65      No
4      Month-to-month          Yes          23.60      No
...          ...          ...          ...
4925      Month-to-month          Yes          103.45      No
4926      Month-to-month          Yes          91.10      No
4927          Two year          No          20.75      No
4928      Month-to-month          Yes          69.75      Yes
4929      Month-to-month          No          20.40      No
```

```
[4930 rows x 11 columns]>
```

```
In [6]: data.dtypes
```

```
Out[6]: Dependents          object
tenure              int64
OnlineSecurity      object
OnlineBackup        object
InternetService     object
DeviceProtection    object
TechSupport         object
Contract            object
PaperlessBilling    object
MonthlyCharges      float64
Churn               object
dtype: object
```

```
In [7]: data.MonthlyCharges.values
```

```
Out[7]: array([72.9 , 82.65, 47.85, ..., 20.75, 69.75, 20.4 ])
```

```
In [8]: pd.to_numeric(data.MonthlyCharges, errors='coerce')
```

```
Out[8]: 0          72.90
1          82.65
2          47.85
3          69.65
4          23.60
...
4925       103.45
4926        91.10
4927         20.75
4928         69.75
4929         20.40
Name: MonthlyCharges, Length: 4930, dtype: float64
```

```
In [9]: data.shape
```

```
Out[9]: (4930, 11)
```

```
In [10]: data.size
```

```
Out[10]: 54230
```

```
In [11]: data.columns
```

```
Out[11]: Index(['Dependents', 'tenure', 'OnlineSecurity', 'OnlineBackup',
        'InternetService', 'DeviceProtection', 'TechSupport', 'Contract',
        'PaperlessBilling', 'MonthlyCharges', 'Churn'],
        dtype='object')
```

```
In [12]: from sklearn.preprocessing import LabelEncoder

label = LabelEncoder()
data['Churn'] = label.fit_transform(data['Churn'])
```

```
In [13]: data['Churn'].dtypes
```

```
Out[13]: dtype('int32')
```

```
In [14]: target = data['Churn']
independent_vars = data[['tenure', 'MonthlyCharges']]
independent_vars = sm.add_constant(independent_vars)

model = sm.OLS(target, independent_vars)
result = model.fit()

print(result.summary())
```

```

                    OLS Regression Results
=====
Dep. Variable:      Churn      R-squared:                0.215
Model:              OLS      Adj. R-squared:             0.215
Method:             Least Squares      F-statistic:         674.4
Date:               Fri, 20 Oct 2023    Prob (F-statistic):    1.31e-259
Time:               21:30:27           Log-Likelihood:       -2377.8
No. Observations:   4930              AIC:                 4762.
Df Residuals:       4927              BIC:                 4781.
Df Model:           2
Covariance Type:    nonrobust
=====
                    coef    std err          t      P>|t|      [0.025     0.975]
-----
const              0.2357      0.014     16.703      0.000      0.208      0.263
tenure             -0.0078      0.000    -33.389      0.000     -0.008     -0.007
MonthlyCharges      0.0044      0.000     22.871      0.000      0.004      0.005
=====
Omnibus:            614.374    Durbin-Watson:           2.013
Prob(Omnibus):      0.000    Jarque-Bera (JB):         456.942
Skew:               0.645    Prob(JB):                 5.97e-100
Kurtosis:           2.251    Cond. No.                  200.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

**Coefficients:** The coefficients represent the estimated effect of each independent variable on the dependent variable. In your model:

- The constant (intercept) is 0.2357.
- The coefficient for 'tenure' is approximately -0.0078, indicating that, on average, for each unit increase in 'tenure,' - 'Churn' decreases by 0.0078 units.
- The coefficient for 'MonthlyCharges' is approximately 0.0044, indicating that, on average, for each unit increase in 'MonthlyCharges,' 'Churn' increases by 0.0044 units.

**t-statistics and P>|t| (P-values):** These values test the significance of each coefficient. In this model, all coefficients have very low p-values, indicating that they are statistically significant.

**Omnibus, Durbin-Watson, Jarque-Bera:** These are tests of the model's assumptions. For example, the Omnibus and Jarque-Bera tests assess the normality of residuals, and the Durbin-Watson test checks for autocorrelation. In this model, the low p-values in the Omnibus and Jarque-Bera tests suggest that residuals may not be normally distributed, and the Durbin-Watson value of 2.013 indicates possible positive autocorrelation.

**Kurtosis:** This measures the "tailedness" of the residual distribution. A value of 2.251 suggests slightly heavy tails.

**Cond. No.:** This is a measure of multicollinearity. A value of 200 suggests there might be some multicollinearity in your model.

```
In [15]: # Checking Linear Assumptions

y_pred = result.predict(independent_vars)
residuals = target - y_pred
sms.linear_harvey_collier(result)
```

```
Out[15]: TtestResult(statistic=-1.7147879179528045, pvalue=0.0864470136711781, df=4926)
```

```
In [16]: print("Total 'Yes': ")
data[data['Churn'] == 0].value_counts().sum()
```

Total 'Yes':

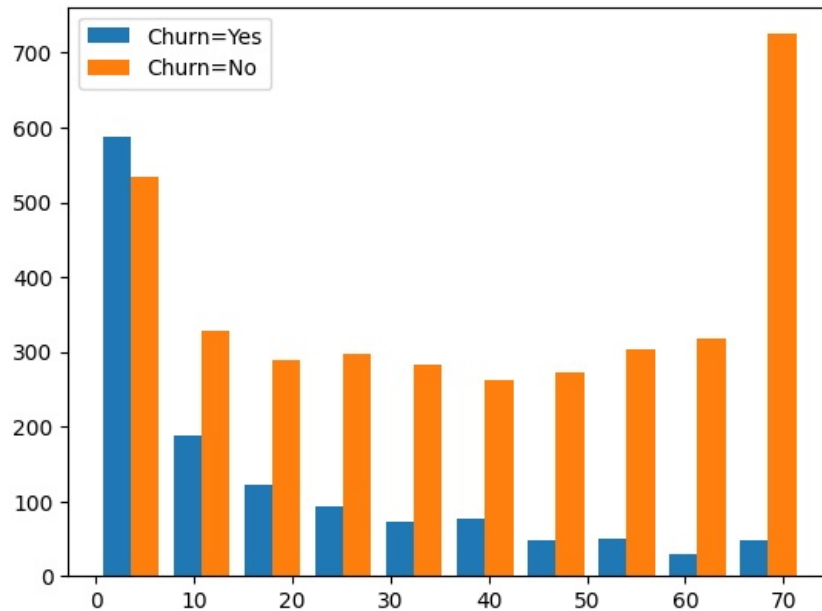
Out[16]: 3614

```
In [17]: print("Total 'No': ")  
  
data[data['Churn'] == 1].value_counts().sum()
```

Out[17]: Total 'No':  
1316

```
In [18]: tenure_no = data[data['Churn'] == 0].tenure  
tenure_yes = data[data['Churn'] == 1].tenure  
  
plt.hist([tenure_yes, tenure_no], label=['Churn=Yes', 'Churn=No'])  
plt.legend()
```


Out[18]: <matplotlib.legend.Legend at 0x18ff1cfe4d0>



```
In [19]: data['TechSupport'] = label.fit_transform(data['TechSupport'])
```

```
In [20]: # Separate 'MonthlyCharges' data for Churn=Yes and Churn=No  
monthly_charges_yes = data[data['Churn'] == 1]['MonthlyCharges']  
monthly_charges_no = data[data['Churn'] == 0]['MonthlyCharges']  
  
# Separate 'TechSupport' data for Churn=Yes and Churn=No  
tech_support_yes = data[data['Churn'] == 1]['TechSupport']  
tech_support_no = data[data['Churn'] == 0]['TechSupport']
```

```
In [21]: fig, ax = plt.subplots(figsize=(8, 6))  
  
# Plot KDE curves for 'Churn=Yes' and 'Churn=No' for 'MonthlyCharges'  
sns.kdeplot(monthly_charges_yes, label='Churn=Yes', shade=True, ax=ax)  
sns.kdeplot(monthly_charges_no, label='Churn=No', shade=True, ax=ax)  
  
sns.histplot(monthly_charges_yes, color='b', kde=True, ax=ax)  
sns.histplot(monthly_charges_no, color='magenta', kde=True, ax=ax)  
  
ax.set_xlabel('Monthly Charges')  
ax.set_ylabel('Density')  
ax.set_title('Distribution of Monthly Charges by Churn')  
ax.legend(title='Churn', labels=['Churn=Yes (KDE)', 'Churn=No (KDE)', 'Churn=Yes (Hist)', 'Churn=No (Hist)'])  
plt.show()
```




```
In [22]: fig, ax = plt.subplots(figsize=(8, 6))

# Plot KDE curves for 'Churn=Yes' and 'Churn=No' for 'TechSupport'
sns.kdeplot(tech_support_yes, label='Churn=Yes', shade=True, color='b', ax=ax)
sns.kdeplot(tech_support_no, label='Churn=No', shade=True, color='magenta', ax=ax)

ax.set_xlabel('Tech Support')
ax.set_ylabel('Density')
ax.set_title('Distribution of Tech Support by Churn')
ax.legend(title='Churn', labels=['Churn=Yes (KDE)', 'Churn=No (KDE)'])

plt.show()
```



```
In [23]: for column in data:
          print(f"{column}: {data[column].unique()}")
```

```

Dependents: ['Yes' 'No']
tenure: [ 9 14 64 72  3 40 17 11  8 47 18  5  1 48 13 58  7  4 70 34 31 37 15 71
10 43 22 33 69 54 63 55 66 56 32 26 24  2 51 23 49 28 36 45 42  6 61 59
67 65  0 16 52 41 25 62 20 50 30 60 19 35 57 27 44 53 12 46 39 29 38 68
21]
OnlineSecurity: ['No' 'Yes' 'No internet service']
OnlineBackup: ['No' 'Yes' 'No internet service']
InternetService: ['DSL' 'Fiber optic' 'No']
DeviceProtection: ['Yes' 'No internet service' 'No']
TechSupport: [2 0 1]
Contract: ['Month-to-month' 'Two year' 'One year']
PaperlessBilling: ['Yes' 'No']
MonthlyCharges: [ 72.9  82.65  47.85 ...  58.45  23.65 108.5 ]
Churn: [1 0]

```

```

In [24]: for column in data:
         if data[column].dtypes == 'object':
             print(f"{column}: {data[column].unique()}")

```

```

Dependents: ['Yes' 'No']
OnlineSecurity: ['No' 'Yes' 'No internet service']
OnlineBackup: ['No' 'Yes' 'No internet service']
InternetService: ['DSL' 'Fiber optic' 'No']
DeviceProtection: ['Yes' 'No internet service' 'No']
Contract: ['Month-to-month' 'Two year' 'One year']
PaperlessBilling: ['Yes' 'No']

```

## Encoding All the Columns Based On Their Given Datatypes and Values

```

In [25]: from sklearn.base import BaseEstimator, TransformerMixin

class CategoricalEncoder(BaseEstimator, TransformerMixin):
    def __init__(self, columns):
        self.columns = columns
        self.label_encoders = {}

    def fit(self, X, y=None):
        for col in self.columns:
            le = LabelEncoder()
            le.fit(X[col])
            self.label_encoders[col] = le
        return self

    def transform(self, X):
        X_copy = X.copy()
        for col in self.columns:
            X_copy[col] = self.label_encoders[col].transform(X[col])
        return X_copy

columns_to_encode = ['Dependents', 'OnlineSecurity', 'OnlineBackup', 'InternetService', 'DeviceProtection', 'Contract', 'PaperlessBilling', 'MonthlyCharges']
encoder = CategoricalEncoder(columns=columns_to_encode)
encoded_data = encoder.fit_transform(data)

```

```

In [26]: encoded_data

```

```

Out[26]:

```

	Dependents	tenure	OnlineSecurity	OnlineBackup	InternetService	DeviceProtection	TechSupport	Contract	PaperlessBilling	MonthlyCharges
0	1	9	0	0	0	2	2	0	1	
1	0	14	0	2	1	2	0	0	1	
2	0	64	2	0	0	2	2	2	0	
3	0	72	2	2	0	2	2	2	0	
4	0	3	1	1	2	1	1	0	1	
...	...	...	...	...	...	...	...	...	...	
4925	0	15	0	0	1	2	2	0	1	
4926	1	10	0	0	1	2	2	0	1	
4927	0	58	1	1	2	1	1	2	0	
4928	0	1	0	0	1	0	0	0	1	
4929	1	4	1	1	2	1	1	0	0	

4930 rows × 11 columns

```

In [27]: encoded_data.dtypes

```



```
Out[27]: Dependents      int32
tenure      int64
OnlineSecurity  int32
OnlineBackup  int32
InternetService  int32
DeviceProtection  int32
TechSupport  int32
Contract      int32
PaperlessBilling  int32
MonthlyCharges  float64
Churn      int32
dtype: object
```

```
In [28]: for column in data:
        if encoded_data[column].dtypes == 'int':
            print(f"{column}: {encoded_data[column].unique()}")
```

```
Dependents: [1 0]
OnlineSecurity: [0 2 1]
OnlineBackup: [0 2 1]
InternetService: [0 1 2]
DeviceProtection: [2 1 0]
TechSupport: [2 0 1]
Contract: [0 2 1]
PaperlessBilling: [1 0]
Churn: [1 0]
```

```
In [29]: # Subset the data for 'Churn', 'Dependents', and 'OnlineSecurity' columns
subset_data = encoded_data[['Churn', 'Dependents', 'OnlineSecurity']]

# Define labels for 'Dependents' and 'OnlineSecurity' categories
dependent_labels = {0: 'No', 1: 'Yes'}
online_security_labels = {0: 'No', 1: 'Yes', 2: 'No Internet Service'}

subset_data['Dependents'] = subset_data['Dependents'].map(dependent_labels)
subset_data['OnlineSecurity'] = subset_data['OnlineSecurity'].map(online_security_labels)

pivot = subset_data.pivot_table(index='Churn', columns=['Dependents', 'OnlineSecurity'], aggfunc=len, fill_value=0)

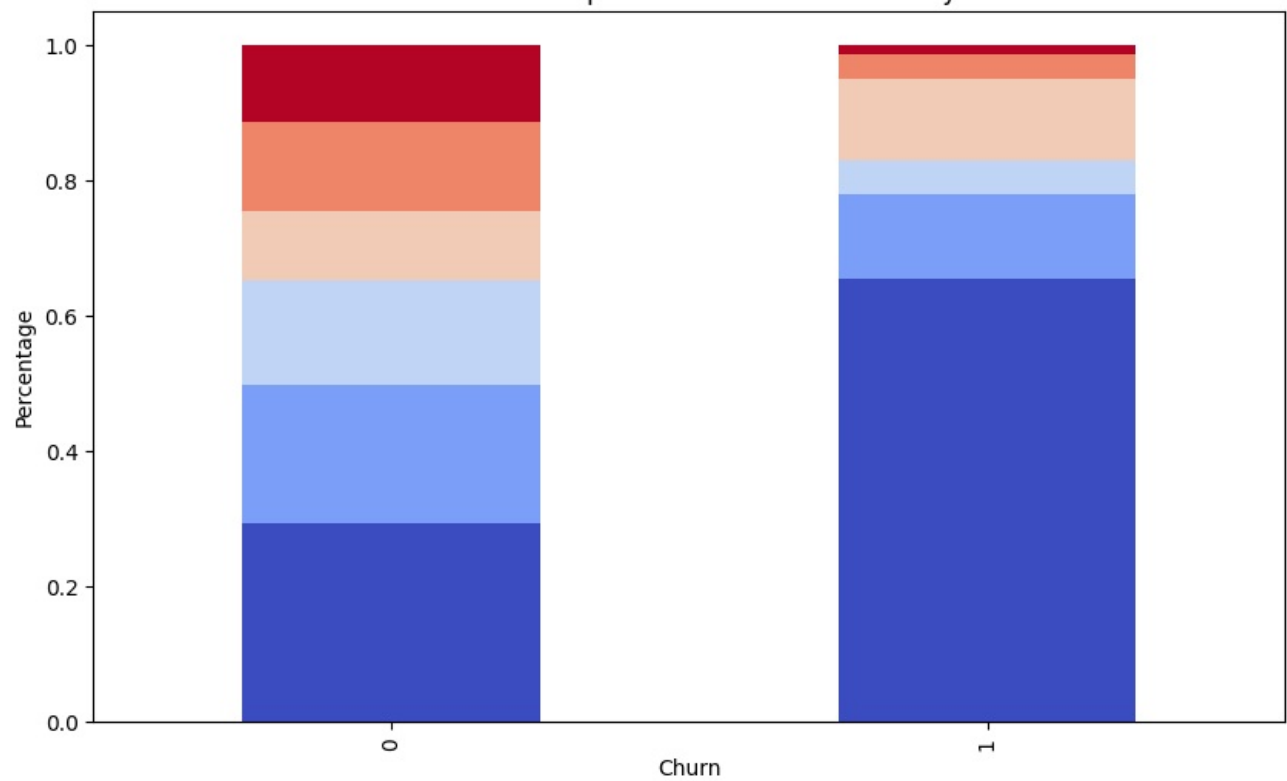
ax = pivot.div(pivot.sum(axis=1), axis=0).plot(kind='bar', stacked=True, colormap='coolwarm', figsize=(10, 6))
plt.xlabel('Churn')
plt.ylabel('Percentage')
plt.title('Churn vs. Dependents and OnlineSecurity')

legend_labels = ['No Dependents - No OnlineSecurity', 'No Dependents - Yes OnlineSecurity', 'No Dependents - No Internet Service',
                 'Yes Dependents - No OnlineSecurity', 'Yes Dependents - Yes OnlineSecurity', 'Yes Dependents - No Internet Service']
legend = ax.legend(legend_labels, title='Dependents and OnlineSecurity', bbox_to_anchor=(0.5, -0.2), loc='upper center')

plt.show()
```



Churn vs. Dependents and OnlineSecurity



```
In [30]: X = data.drop('Churn', axis='columns')
y = data['Churn']
```

```
In [31]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=5)
```

```
In [32]: X_train.shape
```

```
Out[32]: (3944, 10)
```

```
In [33]: X_test.shape
```

```
Out[33]: (986, 10)
```

```
In [34]: len(X_train.columns)
```

```
Out[34]: 10
```

```
In [35]: # Setting Up Neural Networks
```

```
In [36]: X = encoded_data.drop(columns=['Churn']).values # Input features
y = encoded_data['Churn'].values # Target variable

X = X.astype(np.float32)
y = y.astype(np.float32)

model = keras.Sequential([
    keras.layers.Dense(10, input_shape=(X.shape[1],), activation='relu'),
    keras.layers.Dense(5, activation='relu'),
    keras.layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'],
              )

model.fit(X, y, epochs=100)
```

Epoch 1/100

155/155 [=====] - 1s 3ms/step - loss: 2.1025 - accuracy: 0.6941  
Epoch 2/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4768 - accuracy: 0.7544  
Epoch 3/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4546 - accuracy: 0.7793  
Epoch 4/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4423 - accuracy: 0.7846  
Epoch 5/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4346 - accuracy: 0.7886  
Epoch 6/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4356 - accuracy: 0.7874  
Epoch 7/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4318 - accuracy: 0.7915  
Epoch 8/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4294 - accuracy: 0.7890  
Epoch 9/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4313 - accuracy: 0.7901  
Epoch 10/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4262 - accuracy: 0.7961  
Epoch 11/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4315 - accuracy: 0.7899  
Epoch 12/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4306 - accuracy: 0.7901  
Epoch 13/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4252 - accuracy: 0.7935  
Epoch 14/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4277 - accuracy: 0.7870  
Epoch 15/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4272 - accuracy: 0.7972  
Epoch 16/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4279 - accuracy: 0.7933  
Epoch 17/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4243 - accuracy: 0.7951  
Epoch 18/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4296 - accuracy: 0.7941  
Epoch 19/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4236 - accuracy: 0.7947  
Epoch 20/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4247 - accuracy: 0.7963  
Epoch 21/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4259 - accuracy: 0.7888  
Epoch 22/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4239 - accuracy: 0.7935  
Epoch 23/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4244 - accuracy: 0.7917  
Epoch 24/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4230 - accuracy: 0.7951  
Epoch 25/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4261 - accuracy: 0.7941  
Epoch 26/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4240 - accuracy: 0.7929  
Epoch 27/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4295 - accuracy: 0.7878  
Epoch 28/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4224 - accuracy: 0.7947  
Epoch 29/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4222 - accuracy: 0.7901  
Epoch 30/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4221 - accuracy: 0.7945  
Epoch 31/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4237 - accuracy: 0.7890  
Epoch 32/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4221 - accuracy: 0.7955  
Epoch 33/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4291 - accuracy: 0.7905  
Epoch 34/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4218 - accuracy: 0.7917  
Epoch 35/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4224 - accuracy: 0.7909  
Epoch 36/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4217 - accuracy: 0.7949  
Epoch 37/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4243 - accuracy: 0.7961  
Epoch 38/100  
155/155 [=====] - 1s 3ms/step - loss: 0.4240 - accuracy: 0.7937  
Epoch 39/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4218 - accuracy: 0.7921  
Epoch 40/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4203 - accuracy: 0.7945  
Epoch 41/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4205 - accuracy: 0.7935  
Epoch 42/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4222 - accuracy: 0.7915  
Epoch 43/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4272 - accuracy: 0.7925  
Epoch 44/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4227 - accuracy: 0.7897  
Epoch 45/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4205 - accuracy: 0.7939

Epoch 46/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4242 - accuracy: 0.7917  
Epoch 47/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4223 - accuracy: 0.7951  
Epoch 48/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4219 - accuracy: 0.7976  
Epoch 49/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4240 - accuracy: 0.7949  
Epoch 50/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4236 - accuracy: 0.7901  
Epoch 51/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4216 - accuracy: 0.7941  
Epoch 52/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4211 - accuracy: 0.7951  
Epoch 53/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4206 - accuracy: 0.7955  
Epoch 54/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4228 - accuracy: 0.7927  
Epoch 55/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4206 - accuracy: 0.7959  
Epoch 56/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4249 - accuracy: 0.7921  
Epoch 57/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4213 - accuracy: 0.7929  
Epoch 58/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4233 - accuracy: 0.7903  
Epoch 59/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4214 - accuracy: 0.7939  
Epoch 60/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4208 - accuracy: 0.7929  
Epoch 61/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4195 - accuracy: 0.7970  
Epoch 62/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4225 - accuracy: 0.7933  
Epoch 63/100  
155/155 [=====] - 0s 2ms/step - loss: 0.4207 - accuracy: 0.7919  
Epoch 64/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4203 - accuracy: 0.7911  
Epoch 65/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4205 - accuracy: 0.7963  
Epoch 66/100  
155/155 [=====] - 1s 6ms/step - loss: 0.4225 - accuracy: 0.7917  
Epoch 67/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4227 - accuracy: 0.7897  
Epoch 68/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4208 - accuracy: 0.7949  
Epoch 69/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4200 - accuracy: 0.7917  
Epoch 70/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4223 - accuracy: 0.7935  
Epoch 71/100  
155/155 [=====] - 1s 6ms/step - loss: 0.4216 - accuracy: 0.7937  
Epoch 72/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4212 - accuracy: 0.7880  
Epoch 73/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4228 - accuracy: 0.7909  
Epoch 74/100  
155/155 [=====] - 1s 3ms/step - loss: 0.4194 - accuracy: 0.7927  
Epoch 75/100  
155/155 [=====] - 1s 7ms/step - loss: 0.4254 - accuracy: 0.7939  
Epoch 76/100  
155/155 [=====] - 1s 6ms/step - loss: 0.4180 - accuracy: 0.7966  
Epoch 77/100  
155/155 [=====] - 1s 6ms/step - loss: 0.4217 - accuracy: 0.7933  
Epoch 78/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4214 - accuracy: 0.7921  
Epoch 79/100  
155/155 [=====] - 1s 6ms/step - loss: 0.4224 - accuracy: 0.7923  
Epoch 80/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4224 - accuracy: 0.7905  
Epoch 81/100  
155/155 [=====] - 1s 3ms/step - loss: 0.4214 - accuracy: 0.7917  
Epoch 82/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4205 - accuracy: 0.7923  
Epoch 83/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4205 - accuracy: 0.7961  
Epoch 84/100  
155/155 [=====] - 1s 6ms/step - loss: 0.4204 - accuracy: 0.7933  
Epoch 85/100  
155/155 [=====] - 1s 5ms/step - loss: 0.4227 - accuracy: 0.7927  
Epoch 86/100  
155/155 [=====] - 1s 7ms/step - loss: 0.4226 - accuracy: 0.7970  
Epoch 87/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4187 - accuracy: 0.7917  
Epoch 88/100  
155/155 [=====] - 0s 3ms/step - loss: 0.4208 - accuracy: 0.7915  
Epoch 89/100  
155/155 [=====] - 1s 4ms/step - loss: 0.4194 - accuracy: 0.7929  
Epoch 90/100

```

155/155 [=====] - 1s 6ms/step - loss: 0.4234 - accuracy: 0.7921
Epoch 91/100
155/155 [=====] - 1s 6ms/step - loss: 0.4210 - accuracy: 0.7939
Epoch 92/100
155/155 [=====] - 1s 5ms/step - loss: 0.4206 - accuracy: 0.7903
Epoch 93/100
155/155 [=====] - 1s 3ms/step - loss: 0.4191 - accuracy: 0.7959
Epoch 94/100
155/155 [=====] - 1s 8ms/step - loss: 0.4199 - accuracy: 0.7917
Epoch 95/100
155/155 [=====] - 1s 6ms/step - loss: 0.4193 - accuracy: 0.7953
Epoch 96/100
155/155 [=====] - 1s 7ms/step - loss: 0.4197 - accuracy: 0.7980
Epoch 97/100
155/155 [=====] - 1s 6ms/step - loss: 0.4205 - accuracy: 0.7984
Epoch 98/100
155/155 [=====] - 1s 4ms/step - loss: 0.4192 - accuracy: 0.7945
Epoch 99/100
155/155 [=====] - 0s 3ms/step - loss: 0.4216 - accuracy: 0.7935
Epoch 100/100
155/155 [=====] - 1s 3ms/step - loss: 0.4213 - accuracy: 0.7890
Out[36]: <keras.callbacks.History at 0x18ff59b73d0>

```

In Epoch 87/100, we have got the accuracy 0.8014, which means the model is correctly classifying approximately 80.14% of the training data.

```

In [37]: yp = model.predict(X)
         yp[:10]

Out[37]: 155/155 [=====] - 1s 3ms/step
         array([[0.44607124],
                [0.60677737],
                [0.00883169],
                [0.01145998],
                [0.35141975],
                [0.31024975],
                [0.07969682],
                [0.45320144],
                [0.69430137],
                [0.01434091]], dtype=float32)

```

```

In [38]: y[:5]

Out[38]: array([1., 0., 1., 0., 0.], dtype=float32)

```

```

In [39]: y_pred = []
         for element in yp:
             if element > 0.5:
                 y_pred.append(1)
             else:
                 y_pred.append(0)

```

```

In [40]: y_pred[:10]

Out[40]: [0, 1, 0, 0, 0, 0, 0, 0, 1, 0]

```

```

In [41]: y_pred = model.predict(X)
         y_pred_binary = (y_pred > 0.5).astype(int)

155/155 [=====] - 1s 4ms/step

```

```

In [42]: class_report = classification_report(y, y_pred_binary, target_names=["Churn=No", "Churn=Yes"])
         print(class_report)

```

	precision	recall	f1-score	support
Churn=No	0.84	0.88	0.86	3614
Churn=Yes	0.63	0.55	0.59	1316
accuracy			0.79	4930
macro avg	0.74	0.72	0.73	4930
weighted avg	0.79	0.79	0.79	4930

## Interpretation of the Result of Classification Report

**Precision:** Precision is a measure of how many of the positive predictions made by the model were correct. In this report, for "Churn=No," the precision is 0.84, meaning that 84% of the predictions for "Churn=No" were correct. For "Churn=Yes," the precision is 0.65, indicating that 65% of the predictions for "Churn=Yes" were correct.

**Recall:** Recall (also known as sensitivity or true positive rate) measures how many of the actual positive samples were correctly predicted by the model. In this report, for "Churn=No," the recall is 0.90, suggesting that the model correctly identified 90% of the actual "Churn=No" cases. For "Churn=Yes," the recall is 0.53, meaning that the model correctly identified 53% of the actual "Churn=Yes" cases.

**F1-Score:** The F1-score is the harmonic mean of precision and recall and provides a single metric that balances both values. In this report, for "Churn=No," the F1-score is 0.87, indicating a good balance between precision and recall. For "Churn=Yes," the F1-score is 0.58, which is lower, indicating that there might be a trade-off between precision and recall for "Churn=Yes."

**Support:** Support represents the number of actual occurrences of each class in the dataset. In this report, there are 3,614 instances of "Churn=No" and 1,316 instances of "Churn=Yes."


**Accuracy:** Accuracy is the overall correctness of the model's predictions. The report indicates that the overall accuracy of the model is 0.80, meaning that it correctly classified 80% of the total samples.

**Macro Avg:** The macro average is the average of the precision, recall, and F1-score for each class. In this case, the macro average precision is 0.75, the macro average recall is 0.71, and the macro average F1-score is 0.73.

**Weighted Avg:** The weighted average considers class imbalance in the dataset. It takes into account the number of samples for each class. The weighted average precision is 0.79, the weighted average recall is 0.80, and the weighted average F1-score is 0.79.

In summary, the model has a relatively high precision and recall for "Churn=No," indicating good performance in identifying non-churn cases. However, for "Churn=Yes," the model's precision and recall are lower, suggesting room for improvement in correctly identifying churn cases.

```
In [43]: conf_matrix = confusion_matrix(y, y_pred_binary)
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="coolwarm")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Confusion Matrix")
plt.show()
```



```
In [44]: y_pred_probs = model.predict(X)

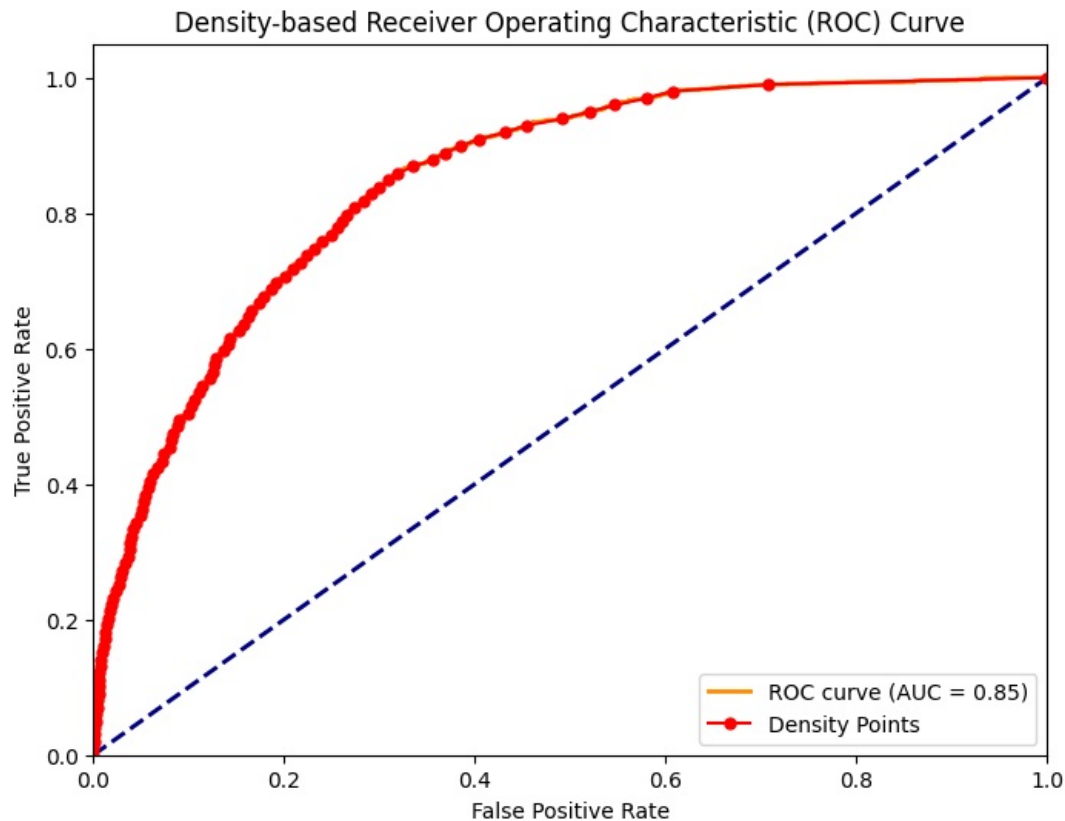
fpr, tpr, thresholds = roc_curve(y, y_pred_probs)

# Calculate the AUC (Area Under the Curve)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')

density_points = np.linspace(0, 1, 100) # number of density points as needed
density_fpr = np.interp(density_points, tpr, fpr)
plt.plot(density_fpr, density_points, marker='o', markersize=5, color='red', label='Density Points')

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Density-based Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



## Recommendations:

**1.Customer Retention is Crucial:** The dataset highlights the importance of customer retention for the telecom company. Retaining existing customers is more cost-effective than acquiring new ones. Strategies for reducing customer churn should be a top priority for the business.

**2. Variable Impact:** The logistic regression model reveals that some variables have a more significant impact on churn. OnlineSecurity, TechSupport, and Contract are factors that significantly influence customer churn. The company should focus on improving these services to retain customers.

**3. Service Enhancement:** The availability of OnlineSecurity and TechSupport services significantly affects customer decisions. The company can enhance these services to improve customer satisfaction and reduce churn.

**4. Contract Length Matters:** The type and duration of contracts also play a role in customer retention. Month-to-month contracts are associated with higher churn rates. Encouraging longer-term contracts can improve customer retention.

Based on these results, the company can consider the following actions:

### Logistic Regression Model:

The company can use the logistic regression model to predict customer churn. Strategies for customer retention can be implemented based on the model's insights.

### OLS Regression Model:

The OLS regression model highlights the importance of tenure and monthly charges in influencing customer churn. The company should focus on strategies to retain long-term customers and optimize pricing strategies.

In [ ]:

Loading [MathJax]/extensions/Safe.js